



An (extremely brief) introduction to Bayesian Markov chain Monte Carlo

Matt Denwood

md@sund.ku.dk

Department of Veterinary and Animal Sciences

University of Copenhagen

Denmark



Overview

- What is Bayesian statistics?
- What is the theory behind MCMC?
- How does MCMC work in practice?
- A basic example: prevalence estimation



What is Bayesian statistics?



An example using diagnostic tests

A well-known formula in epidemiology:

$$PPV = \frac{TP}{TP + FP} = \frac{Se \cdot Prev}{Se \cdot Prev + (1 - Sp) \cdot (1 - Prev)}$$

But where does this come from?

$$Se = P(T^+ | D^+)$$

$$Sp = P(T^- | D^-) = 1 - P(T^+ | D^-)$$

$$PPV = P(D^+ | T^+)$$

$$Prev = P(D^+) = 1 - P(D^-)$$

Likelihood

Prior information

$$P(D^+ | T^+) = \frac{P(T^+ | D^+) \cdot P(D^+)}{P(T^+ | D^+) \cdot P(D^+) + P(T^+ | D^-) \cdot P(D^-)}$$

Posterior information



Bayes theorem

Discrete form:

$$P(A_j | B) = \frac{P(B | A_j) \cdot P(A_j)}{\sum_i P(B | A_i) \cdot P(A_i)}$$

Continuous form:




$$f_X(x | Y = y) = \frac{f_Y(y | X = x) \cdot f_X(x)}{\int_{-\infty}^{\infty} f_Y(y | X = \varepsilon) \cdot f_X(\varepsilon) \cdot d\varepsilon}$$

General form:

Posterior \propto Likelihood \times Prior

Bayesian statistics

- "Belief" is expressed as a probability distribution
 - i.e. We don't know exactly what the parameter value is, but we can quantify our uncertainty
 - This is the key difference to frequentist statistics
- We usually have continuous probability distributions for both prior and posterior



What we knew before + What the data tell us = What we know now

What is the theory behind MCMC?



MCMC: the theory

First define a function to calculate a posterior probability corresponding to a Binomial likelihood and Beta prior:

```
log_posterior_fun <- function(parameter){  
  ... See the exercise if you are interested in R code ...  
}
```

Then choose a 'reasonable' place to start looking for our parameter values, and calculate the posterior at that parameter value:

```
parameter[1] <- 0.25  
log_post[1] <- log_posterior_fun(parameter[1])  
log_post[1]  
-27.5222
```



MCMC: the theory

Now sample another parameter value using a random walk:

```
new_par <- rnorm(1, mean=parameter[1], sd=sigma)
          0.2606076
new_lpost <- log_posterior_fun(new_par)
          -18.00832
```

That's a bit better! Let's store this new value:

```
parameter[2] <- new_par
log_post[2] <- new_lpost
```

And resample:

```
new_par <- rnorm(1, mean=parameter[1], sd=sigma)
          0.2572526
new_lpost <- log_posterior_fun(new_par)
          -20.72975
```



MCMC: the theory

This is a bit worse. Should we keep it or not?

We will let fate decide ... but it makes sense to have a higher chance of accepting the new parameter value if it is not much worse than the old value:

```
probability_ratio <- exp(new_lpost - log_post[2])  
0.06578108  
accept <- rbinom(1, 1, probability_ratio)  
0
```

So we reject this new value, and keep the old one

We also reward the old parameter for being hard to beat by counting it twice (so this counts as the 2nd *and* 3rd values in the chain)

```
parameter[3] <- parameter[2]  
log_post[3] <- log_post[2]
```




MCMC: the theory

Repeat for the rest of the 1000 iterations, then look at the results:

```
plot(parameter, type='l')
```

Trace plot of parameter values



Markov chain Monte Carlo

This algorithm is a basic implementation of MCMC!

- MCMC is used to approximate a posterior distribution by sampling from it
- Dependence on previous value is critical
 - Consecutive samples are *not* independent
 - Eventually a Markov chain will forget where it started, but it might take a while before it does this

This non-randomness of Markov chains has 2 important consequences:

1. The starting values are not immediately forgotten
 - We need a *burn-in* period before sampling from the chains (aka identifying convergence on the stationary distribution)
2. A sample of 1000 auto-correlated iterations gives us less information than a sample of 1000 independent iterations
 - We need to know the *effective sample size* or *Monte Carlo error*




1. Burn-in

It is ESSENTIAL to identify and remove the burn-in period

```
plot(parameter, type='l')
```

Trace plot of parameter values




1. Burn-in

It is ESSENTIAL to identify and remove the burn-in period


```
plot(parameter[-(1:50)], type='l')
```

Trace plot of stationary chain



2. Effective sample size


Histogram of 950 sampled values



An effective sample size of 137 – not a good approximation

2. Effective sample size


Histogram of 10000 sampled values



An effective sample size of 1752 – a pretty good approximation

2. Effective sample size

Histogram of 100k sampled values



An effective sample size of 16759 – a very good approximation

Dangers of MCMC

1. Problems related to the Markov chain
 - We didn't remove enough burn-in iterations so our posterior is skewed by the parameter space from where we happened to start the initial values
 - Our chain hasn't really found the stationary posterior, and is just sampling from a local solution
 - **CONVERGENCE**
2. Problems related to the Monte Carlo integration
 - We didn't take enough samples to adequately represent the true posterior
 - **EFFECTIVE SAMPLE SIZE / MONTE CARLO ERROR**



How does MCMC work in practice?



Using BUGS

Bayesian (analysis) Using Gibbs Sampling

- ✓ Generic MCMC framework with ready-made algorithms
- ❖ Does not automatically resolve any of the dangers with MCMC sampling!

A BUGS model is compiled by the software into an MCMC sampler

- All variables in the model must be defined EXACTLY once
- NB: it is a symbolic language NOT a programming language

OpenBUGS

- Available for Windows, Linux and Mac (under emulation)
- No longer under active development

JAGS (Just Another Gibbs Sampler)

- Technically not BUGS, but almost identical
- Primarily designed to be called from within R
 - Packages available: rjags, runjags, R2jags, jagsUI




Procedure

- 1) Define a syntactically valid model (used to calculate the posterior)
- 2) Supply data and initial values for at least 2 chains (these are often optional)
- 3) Decide which variables to monitor
- 4) Run the model
- 5) Make sure the chains have converged
 - Brooks-Gelman-Rubin statistic (or BGR) compares the variance within chains to the variance between chains – ratio < 1.05 indicates that chains are sampling the same values so have probably converged
 - But always check the trace plots visually as well!
- 6) Make sure the number of sampled iterations is high enough
 - Minimum effective sample size of >400 for parameters of interest
 - Or equivalently when $MC_error < 0.05 \times \text{sample SD}$



Visual assessment of convergence


Have these two chains (pink and blue) converged?



Looks good
[We are OK to look at the model results]

Visual assessment of convergence

Have these two chains (pink and blue) converged?



Looks bad
[Model results cannot be used]

A basic example: prevalence estimation



JAGS: Define and run the model

Define the model, load the runjags library, and define the data in R:

```
bugs_model <- "
  model{
    y ~ dbin(ap, n)
    #Uniform (non-informative) prior for apparent prev.
    ap ~ dbeta(1,1)

    #data# y, n
    #inits# ap
    #monitor# ap
  }
"

library('runjags')

y <- 1210
n <- 4072
ap <- 0.1

results <- run.jags(bugs_model, n.chains=2, burnin=5000,
  sample=10000)
```



OpenBUGS: Define and run the model

The screenshot displays the OpenBUGS application window. The main text area contains the following code:

```
#model specification
model {
y ~ dbin(ap, n)

#Uniform (non-informative) prior for apparent prevalence (ap)
ap ~ dbeta(1,1)
}

#data
list(n=4072, y=1210)





#initials
list(ap=0.1)
```

Three utility windows are open on the right side of the interface:

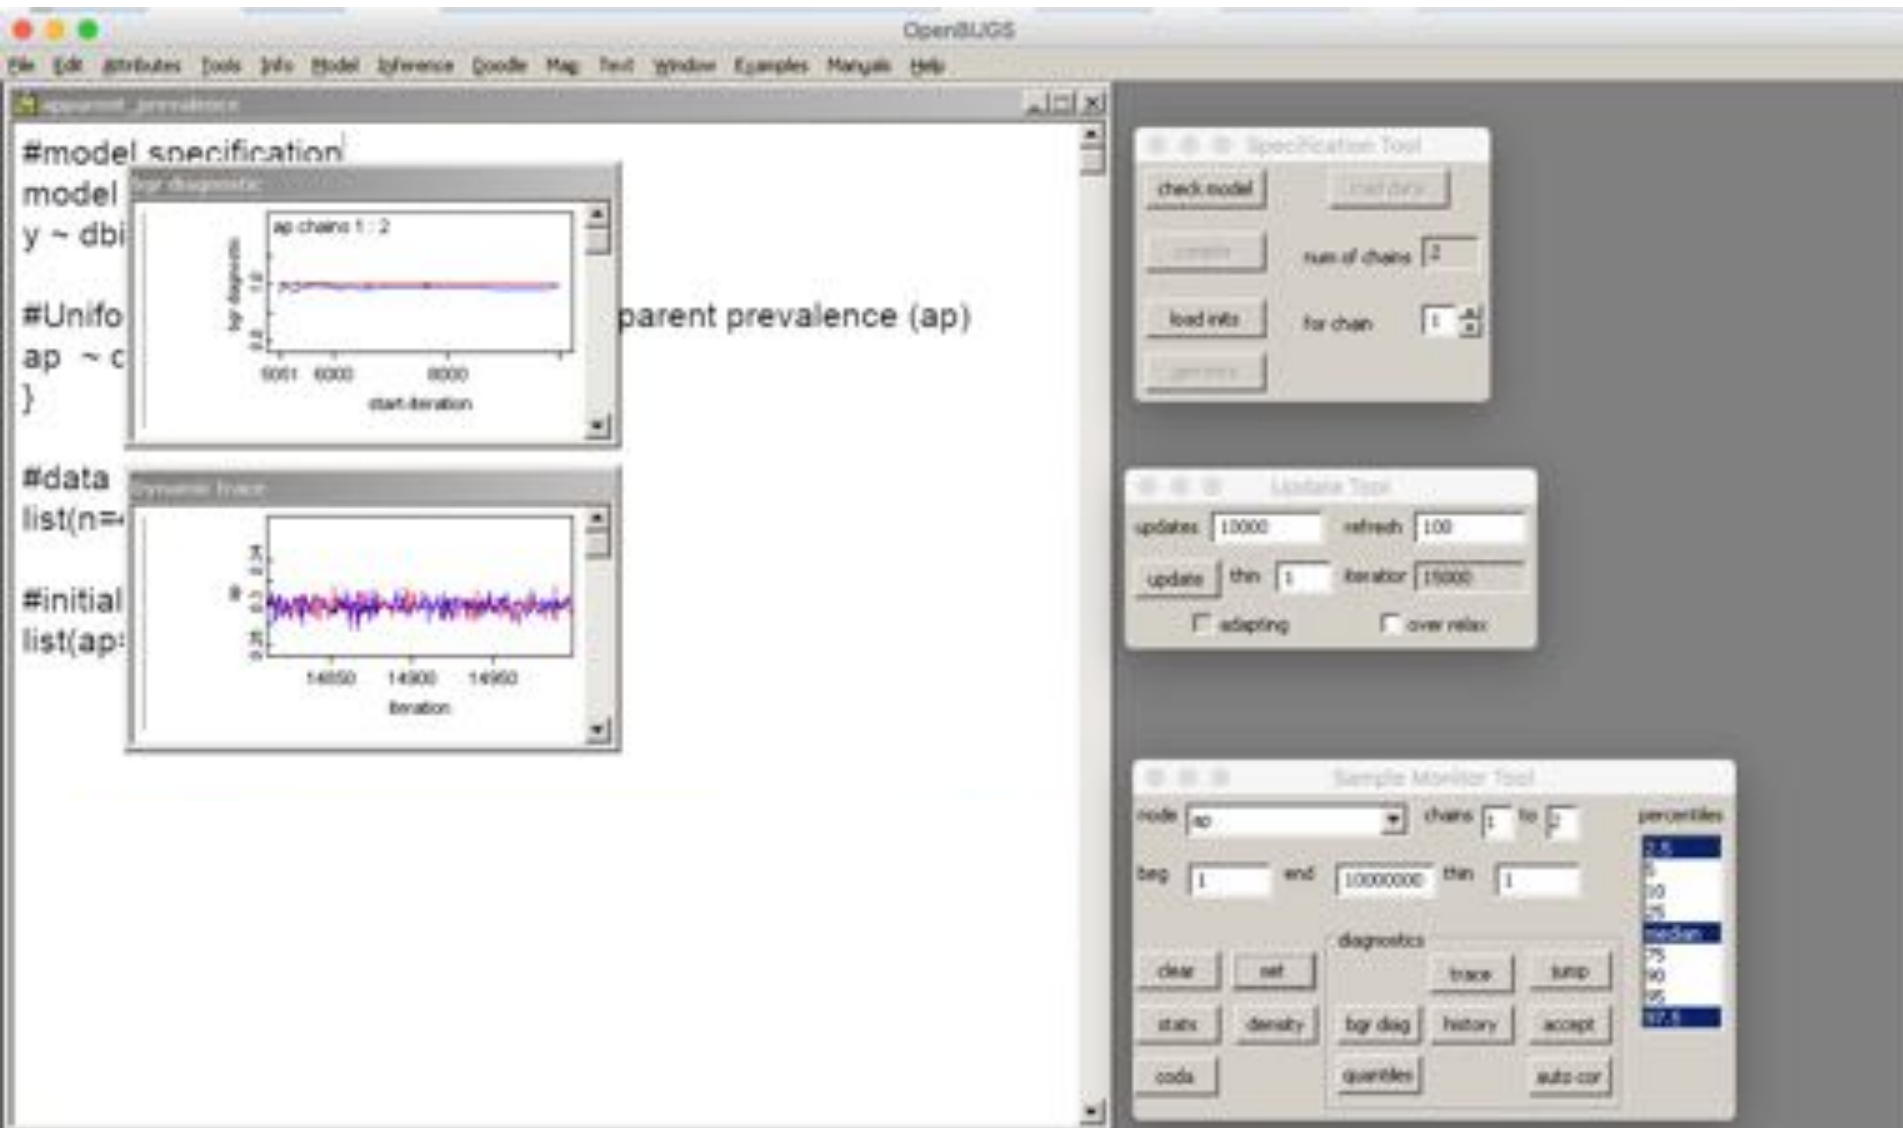
- Specification Tool:** Contains buttons for 'check model', 'load data', 'compile', 'load inits', and 'update inits'. It also has input fields for 'run of chains' (set to 2) and 'for chain' (set to 1).
- Update Tool:** Contains input fields for 'updates' (5000), 'refresh' (100), 'update' (button), 'thin' (1), and 'iteration' (5000). It also has checkboxes for 'adapting' and 'over relax'.
- Sample Monitor Tool:** Contains a 'node' dropdown, 'chain' dropdown (1 to 2), 'beg' (1), 'end' (1000000), and 'thin' (1) fields. It has a 'diagnostics' section with buttons for 'clear', 'plot', 'status', 'density', 'log-plot', 'histogram', 'trace', 'plot', 'auto', and 'plot auto'. On the right, a 'percentiles' list shows values: 2.5, 5, 10, 25, median, 75, 90, 95, 97.5.

JAGS: Check convergence

`plot(results)`



OpenBUGS: Check convergence



JAGS: Check sample size and results

```
> results
```

JAGS model summary statistics from 20000 samples (chains = 2;
adapt+burnin = 6000):

	Lower95	Median	Upper95	Mean	SD	Mode
ap	0.28339	0.29719	0.31155	0.2972	0.0071436	--

	MCerr	MC%ofSD	SSeff	AC.10	psrf
ap	0.000063509	0.9	12652	0.018843	1.0002

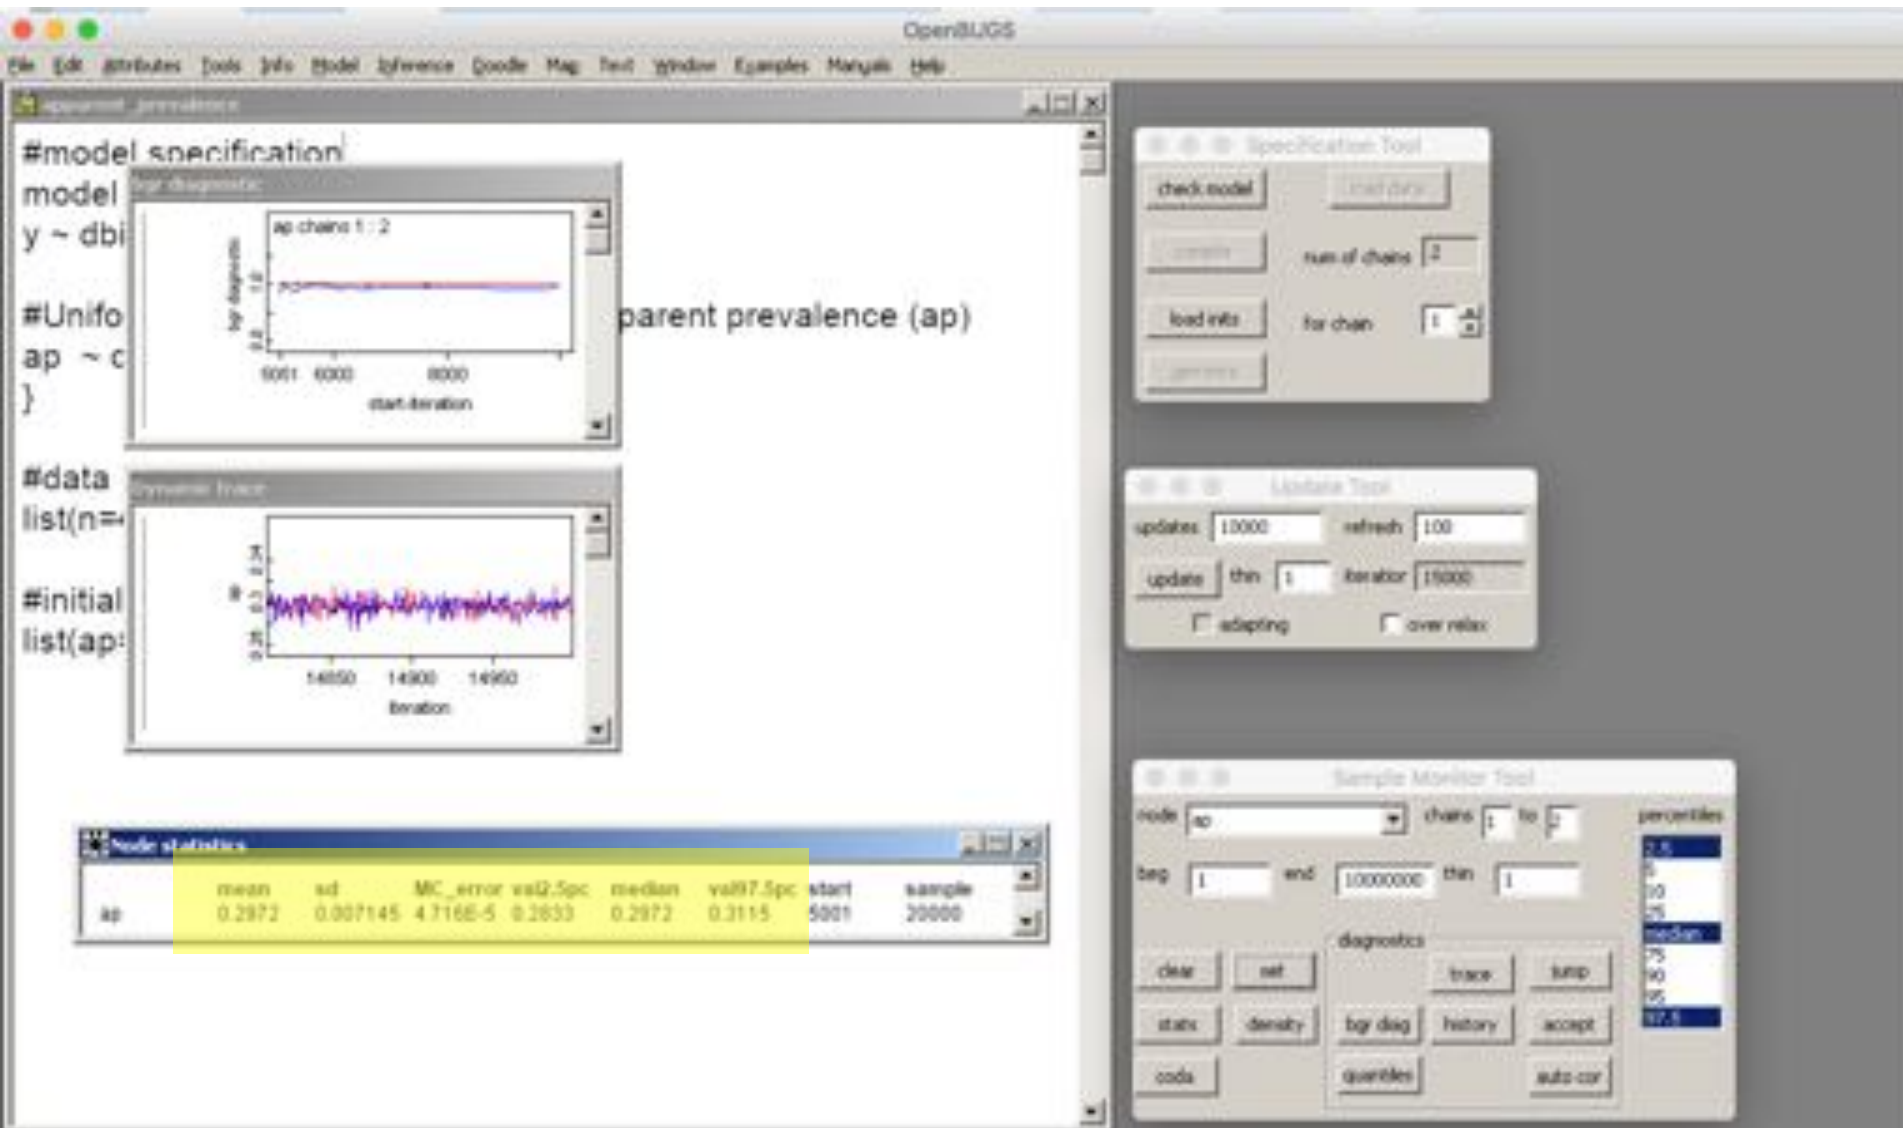
Total time taken: 0.8 seconds

See also:

?runjags



OpenBUGS: Check sample size and results



Practical 1: apparent prevalence

- Decide if you are using JAGS or OpenBUGS
 - JAGS is recommended (by me) unless you really hate R...
- For JAGS
 - Open the apparent_prevalence.R file in R
 - Read and run the code down to line 34
 - Look at the help files as indicated
 - If you are interested (and reasonably proficient in R) check out the Metropolis algorithm from lines 48-124
- For OpenBUGS
 - Open the apparent_prevalence.odc file in OpenBUGS
 - Follow along with group demonstration
- Next presentation: 9.50am

Advert: ABME05 course will run 11th-15th November 2019
<https://www.prstatistics.com/course/applied-bayesian-modelling-for-ecologists-and-epidemiologists-abme05/>